

Databases

Lecture 4

Querying Relational Databases Using SQL (II)

Functional Dependencies. Normal Forms

- **GROUP BY, HAVING**

```
SELECT [DISTINCT] select-list  
FROM from-list  
WHERE qualification  
GROUP BY grouping-list  
HAVING group-qualification
```

- optional GROUP BY clause
 - list of (expressions involving) columns used for grouping
- optional HAVING clause
 - group qualification conditions
- aggregation operators

COUNT, AVG, SUM, MIN, MAX

- **GROUP BY, HAVING**

```
SELECT [DISTINCT] select-list  
FROM from-list  
WHERE qualification  
GROUP BY grouping-list  
HAVING group-qualification
```

- **group**

- a collection of rows with identical values for the columns in **grouping-list**
- every row in the result of the query corresponds to a group

- **GROUP BY, HAVING**

```
SELECT [DISTINCT] select-list  
FROM from-list  
WHERE qualification  
GROUP BY grouping-list  
HAVING group-qualification
```

- **select-list**

- columns (that must appear in **grouping-list**)
- terms of the form *aggop(column) [AS NewName]*
 - e.g., MAX(R.ImpactFactor) AS MaxImpactFactor
 - *NewName* assigns a name to the column in the result table

- **GROUP BY, HAVING**

SELECT [DISTINCT] select-list

FROM from-list

WHERE qualification

GROUP BY grouping-list

HAVING group-qualification

- **group-qualification**

- expressions with a single value / group
- a column in **group-qualification** appears in **grouping-list** or as an argument to an aggregation operator
- records that meet **qualification** are partitioned into groups based on the values of the columns in **grouping-list**
- an answer row is generated for every group that meets **group-qualification**

Find the age of the youngest researcher for each impact factor.

```
SELECT R.ImpactFactor, MIN(R.Age)
FROM Researchers R
GROUP BY R.ImpactFactor
```

* discussion: using the GROUP BY clause vs writing n queries, one for each of the n values of the impact factor, where n depends on the relation instance

Find the age of the youngest researcher who is at least 18 years old for each impact factor with at least 10 such researchers.

```
SELECT R.ImpactFactor, MIN(R.Age) AS MinAge
FROM Researchers R
WHERE R.Age >= 18
GROUP BY R.ImpactFactor
HAVING COUNT(*) >= 10
```

Find the average age of researchers who are at least 18 years old for each impact factor with at least 10 researchers.

```
SELECT R.ImpactFactor, AVG(R.Age) AS AvgAge
FROM Researchers R
WHERE R.Age >= 18
GROUP BY R.ImpactFactor
HAVING 9 < (SELECT COUNT(*)
            FROM Researchers R2
            WHERE R2.ImpactFactor = R.ImpactFactor)
```

Problem

See lecture query: *Find the most prolific researchers, i.e., those with the largest number of papers*, on schema *AuthorContribution*(*RID*, *PID*, *Year*).

Does it compute the same result on *AuthorContribution*(*RID*, *PID*, *Year*)? If not, change it so the intended result is computed on this schema as well.

Find the name and age of the oldest researcher.

```
SELECT R.Name, MAX(R.Age)
```

```
FROM Researchers R
```

--**error**: if the SELECT clause contains an aggregation operator, then it must contain **only** aggregation operators, unless the query has a GROUP BY clause

Correct query:

```
SELECT R.Name, R.Age
```

```
FROM Researchers R
```

```
WHERE R.Age = (SELECT MAX(R2.Age)  
               FROM Researchers R2)
```

- ORDER BY, TOP

Sort researchers by impact factor (in descending order) and age (in ascending order).

```
SELECT *  
FROM Researchers R  
ORDER BY R.ImpactFactor DESC, R.Age ASC
```

Retrieve the names and ages of the top 10 researchers ordered by name.

```
SELECT TOP 10 R.Name, R.Age  
FROM Researchers R  
ORDER BY R.Name
```

Find the top 25% researchers (all the data) ordered by age (descending).

```
SELECT TOP 25 PERCENT *  
FROM Researchers R  
ORDER BY R.Age DESC
```

Find the number of researchers for each impact factor. Order the result by the number of researchers.

```
SELECT R.ImpactFactor, COUNT(*) AS NoR  
FROM Researchers R  
GROUP BY R.ImpactFactor  
ORDER BY NoR
```

- obs. intersection queries can be expressed with IN.

Find the names of researchers who have published in IDEAS and EDBT.

```
SELECT R.Name
FROM Researchers R INNER JOIN AuthorContribution A
    ON R.RID = A.RID
    INNER JOIN Papers P ON A.PID = P.PID
WHERE P.Conference = 'IDEAS' AND
    R.RID IN (SELECT A2.RID
        FROM AuthorContribution A2 INNER JOIN Papers P2
            ON A2.PID = P2.PID
        WHERE P2.Conference = 'EDBT')
```

- obs. set-difference queries can be expressed with NOT IN.

Find the researchers who have published in EDBT but not in IDEAS.

```
SELECT A.RID
FROM AuthorContribution A INNER JOIN Papers P ON A.PID = P.PID
WHERE P.Conference = 'EDBT' AND
      A.RID NOT IN (SELECT A2.RID
                    FROM AuthorContribution A2 INNER JOIN Papers P2
                    ON A2.PID = P2.PID
                    WHERE P2.Conference = 'IDEAS')
```

Find researchers whose IF is greater than the IF of some researcher called *Ionescu*.

```
SELECT R.RID
FROM Researchers R
WHERE R.ImpactFactor >
      (SELECT MIN(R2.ImpactFactor)
       FROM Researchers R2
       WHERE R2.Name = 'Ionescu')
```

Find researchers whose IF is greater than the IF of every researcher called *Ionescu*.

```
SELECT R.RID
FROM Researchers R
WHERE R.ImpactFactor >
      (SELECT MAX(R2.ImpactFactor)
       FROM Researchers R2
       WHERE R2.Name = 'Ionescu')
```

- the SELECT statement:

$$\begin{aligned}
 & \text{SELECT } \left[\left\{ \begin{array}{c} \text{ALL} \\ \text{DISTINCT} \\ \text{TOP } n \text{ [PERCENT]} \end{array} \right\} \right] \left\{ \text{expr1 [AS column1]} [, \text{expr2 [AS column2]}] \dots \right\}^* \\
 & \text{FROM source1 [alias1]} [, \text{source2 [alias2]}] \dots \\
 & [\text{WHERE qualification}] \\
 & [\text{GROUP BY grouping_list}] \\
 & [\text{HAVING group_qualification}] \\
 & \left[\left\{ \begin{array}{c} \text{UNION [ALL]} \\ \text{INTERSECT} \\ \text{EXCEPT} \end{array} \right\} \text{SELECT_statement} \right] \\
 & \left[\text{ORDER BY } \left\{ \begin{array}{c} \text{column1} \\ \text{column1_number} \end{array} \right\} \left[\left\{ \begin{array}{c} \text{ASC} \\ \text{DESC} \end{array} \right\} \right] \left[', \left\{ \begin{array}{c} \text{column2} \\ \text{column2_number} \end{array} \right\} \left[\left\{ \begin{array}{c} \text{ASC} \\ \text{DESC} \end{array} \right\} \right] \right] \dots \right] \right]
 \end{aligned}$$

- non-procedural query
- SELECT statement evaluation: the result is a relation (table)
- data can be obtained from one or multiple data sources; a source can have an associated *alias*, used only in the SELECT statement
- various expressions are evaluated on the data (from the above-mentioned sources)
- a source column can be qualified with the source's name (or alias)

- a data source can be:
 1. table / view in the database
 2. (SELECT_statement)
 3. join_expression:
 - source1 [alias1] join_operator source2 [alias2] ON join_condition
 - (join_expression)
- * a join condition can be of the form:
 - elementary_cond
 - (condition)
 - NOT condition
 - condition1 AND condition2
 - condition1 OR condition2

* an elementary join condition (elementary_cond) among two data sources can be of the form:

- [source1_alias.]column1 relational_operator [source2_alias.]column2
- expression1 relational_operator expression2 (expression1 and expression2 use columns from different sources)

- the WHERE clause can contain filter and join conditions
- filter conditions:
 - expression relational_operator expression
 - expression [NOT] BETWEEN valmin AND valmax
 - expression [NOT] LIKE pattern
 - expression IS [NOT] NULL
 - expression [NOT] IN (value [, value] ...)
 - expression [NOT] IN (subquery)
 - expression relational_operator {ALL | ANY} (subquery)
 - [NOT] EXISTS (subquery)
- filter conditions can be:
 - elementary (described above)
 - composed with logical operators and parentheses

- obs: not all DBMSs support TOP
 - MySQL: `SELECT ... LIMIT n`
 - Oracle: `SELECT ... WHERE ROWNUM <= n`
- rules for building expressions:
 - operands: constants, columns, system functions, user functions
 - operators: corresponding to operands
- ordering records: the `ORDER BY` clause

- the SELECT statement - logical processing (Transact-SQL)

FROM

WHERE

GROUP BY

HAVING

SELECT

DISTINCT

ORDER BY

TOP

Functional Dependencies. Normal Forms


- there are multiple ways to model a data collection in the relational model (using tables)
- database design should be carried out in such a manner that subsequent queries and operations are performed as smoothly as possible, i.e.:
 - no additional tests are required when changing data
 - operations can be performed through SQL statements alone
- the above mentioned operations can be easily carried out when relations satisfy certain conditions (i.e., relations are in a certain normal form)

data redundancy

- generates a series of problems; some can be tackled by replacing a relation with a collection of smaller relations; each of the latter contains a (strict) subset of attributes from the original relation

Appendix - More SQL Queries

- more query examples on the database:



- A student's trajectory (academic year and group).

```
SELECT acad_year, sgroup  
FROM Trajectory  
WHERE program = 2 AND reg_number = '7654'
```

- A student's grades.

```
SELECT acad_year, course, no_credits, grade  
FROM Results  
WHERE program = 2 AND reg_number = '7654'
```

- **Students who belonged to group 915 in the academic year 2017-2018.**

```
SELECT last_name, first_name, s.reg_number
FROM Students s INNER JOIN Trajectory t
    ON s.program=t.program AND s.reg_number=t.reg_number
WHERE acad_year=2017 AND sgroup='915'
```

```
SELECT last_name, first_name, s.reg_number
FROM Students s INNER JOIN
    (SELECT *
     FROM Trajectory
     WHERE acad_year=2017 AND sgroup='915') t
ON s.reg_number=t.reg_number AND s.program=t.program
```

- Students who belong to a group, but have no grades - in the academic year 2017-2018.

```
SELECT last_name, first_name
FROM Students AS s
WHERE EXISTS (SELECT *
              FROM Trajectory t
              WHERE acad_year=2017 AND t.program=s.program
              AND t.reg_number=s.reg_number)
AND NOT EXISTS (SELECT *
                FROM Results r
                WHERE acad_year=2017 AND
                      s.program=r.program AND
                      s.reg_number=r.reg_number)
--ORDER BY last_name, first_name
```

- Students who belong to a group, but have no grades - in the academic year 2017-2018.

```
SELECT last_name, first_name
FROM (Students s INNER JOIN
      (SELECT *
       FROM Trajectory WHERE acad_year=2017) t
      ON s.program=t.program AND s.reg_number=t.reg_number
)
LEFT JOIN
      (SELECT *
       FROM Results
       WHERE acad_year=2017) r
      ON s.program=r.program AND s.reg_number=r.reg_number
WHERE grade IS NULL
```

- The number of students in the database.

```
SELECT COUNT(*) AS NoS  
FROM Students
```

- The number of students born on the same day, regardless of year and month.

```
SELECT SUBSTRING(cnp, 6, 2) AS day, COUNT(*) AS NoS  
FROM Students  
GROUP BY SUBSTRING(cnp, 6, 2)
```

- The grades of a given student (only the maximum grade is required for each course).

```
SELECT course, no_credits, MAX(grade) AS maxgrade  
FROM Results  
WHERE program = 3 AND reg_number='virtualstudent11'  
GROUP BY course, no_credits  
--ORDER BY course
```

- The grades of a given student (only the maximum grade is required for each course). Include the name of the course in the answer set.

```
SELECT id, name, no_credits, maxgrade AS grade
FROM Courses c INNER JOIN
    (SELECT course, no_credits, MAX(grade) AS maxgrade
     FROM Results
     WHERE program = 3 AND reg_number='virtualstudent11'
     GROUP BY course, no_credits) r
ON c.id = r.course
--ORDER BY name
```


- For each student name that appears at least 3 times, retrieve all students with that name.

```
SELECT *  
FROM Students  
WHERE last_name IN  
    (SELECT last_name  
     FROM Students  
     GROUP BY last_name  
     HAVING COUNT(*) >= 3)  
--ORDER BY last_name, first_name
```

*** rewrite the query without IN**

- The number of students in each program and year of study in 2018.

```
SELECT g.program, g.year_of_study, COUNT(*) AS NoS
FROM Trajectory t INNER JOIN Groups g
  ON t.program=g.program AND t.sgroup=g.id AND
   t.acad_year=g.acad_year
WHERE t.acad_year=2018
GROUP BY g.program, g.year_of_study
```

- The last name, first name, program (id), registration number, number of passed exams, number of credits, and gpa for each student with at least 30 credits in passed courses at the end of 2017.

```
SELECT last_name, first_name, s.program, s.reg_number, COUNT(*)
       AS noc, SUM(no_credits) AS nocr,
       SUM(graded*no_credits)/SUM(no_credits) AS gpa
FROM Students s INNER JOIN
     (SELECT program, reg_number, course, no_credits, MAX(grade)
      AS graded
FROM Results
WHERE acad_year=2017 AND grade>=5
GROUP BY program, reg_number, course, no_credits
) r
ON s.program=r.program AND s.reg_number=r.reg_number
GROUP BY s.program, s.reg_number, last_name, first_name
HAVING SUM(no_credits)>=30
ORDER BY 3,1,2
```

- For every course, the number of grades and the number of passing grades in 2017.

*** a. MySQL ***

```
SELECT name, COUNT(*) AS nrg, SUM(IF(r.grade >= 5,1,0))  
AS nrpg  
FROM Courses c INNER JOIN  
    (SELECT * FROM Results WHERE acad_year = 2017) r  
    ON c.id = r.course  
GROUP BY name  
--ORDER BY 1
```

- For every course, the number of grades and the number of passing grades in 2017.

*** b. Oracle ***

```
SELECT name, COUNT(*) AS nrg, SUM(CASE WHEN (r.grade >= 5) THEN
1 ELSE 0 END) AS nrpg
FROM Courses c INNER JOIN
    (SELECT * FROM Results WHERE acad_year = 2017) r
    ON c.id = r.course
GROUP BY name
--ORDER BY 1
```

- * The grade average for each group in 2017. *

a. Maximum & passing grades in 2017.

```
CREATE VIEW studgrades AS
SELECT program, reg_number, course, no_credits, MAX(grade) AS
    mgrade
FROM Results
WHERE acad_year=2017 AND grade>=5
GROUP BY program, reg_number, course, no_credits
```

b. Students' grade average (in 2017).

```
CREATE VIEW studavg AS
SELECT program, reg_number,
    SUM(no_credits*mgrade)/SUM(no_credits) AS gradeavg
FROM studgrades
GROUP BY program, reg_number
HAVING SUM(no_credits) >= 30
```

- * The grade average for each group in 2017. *

c. Final query:

```
SELECT sgroup, AVG(gradeavg) AS gravg
FROM
    (SELECT program, reg_number, sgroup
     FROM Trajectory
     WHERE acad_year=2017) t
     INNER JOIN studavg AS s
         ON t.reg_number=s.reg_number AND t.program=s.program
GROUP BY sgroup
```

- what happens if two groups in different programs have the same id?

- * The grade average for each program in 2017. *

```
CREATE VIEW programsavg AS  
SELECT id, name, AVG(gradeavg) AS pravg  
FROM Programs p INNER JOIN studavg s ON p.id=s.program  
GROUP BY id, name
```

```
SELECT * FROM programsavg
```


References

- [Ta13] ȚÂMBULEA, L., Curs Baze de date, Facultatea de Matematică și Informatică, UBB, 2013-2014
- [Ra02] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems (3rd Edition), McGraw-Hill, 2002
- [Da03] DATE, C.J., An Introduction to Database Systems (8th Edition), Addison-Wesley, 2003
- [Ga09] GARCIA-MOLINA, H., ULLMAN, J., WIDOM, J., Database Systems: The Complete Book (2nd Edition), Pearson Education, 2009
- [Ha96] HANSEN, G., HANSEN, J., Database Management And Design (2nd Edition), Prentice Hall, 1996
- [Ra02S] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems, Slides for the 3rd Edition,
<http://pages.cs.wisc.edu/~dbbook/openAccess/thirdEdition/slides/slides3ed.html>
- [Si11] SILBERSCHATZ, A., KORTH, H., SUDARSHAN, S., Database System Concepts (6th Edition), McGraw-Hill, 2011
- [Si19S] SILBERSCHATZ, A., KORTH, H., SUDARSHAN, S., Database System Concepts, Slides for the 7th Edition, <http://codex.cs.yale.edu/avi/db-book/>
- [Ul11] ULLMAN, J., WIDOM, J., A First Course in Database Systems,
<http://infolab.stanford.edu/~ullman/fcdb.html>